

The Model is not the Moat

Published: AUGUST 3, 2026

Source: <https://www.franklintempleton.com/articles/en/2026/08/the-model-is-not-the-moat>

Extraction method: selenium

Key takeaways

The model is becoming a commodity. The system is becoming the moat. As models become increasingly interchangeable, durable value shifts toward the infrastructure, workflows, governance and proprietary context that turn intelligence into trusted work.

Enterprises don't buy intelligence. They buy outcomes. The relevant measure isn't cost per token—it's the cost of completing a task reliably, securely and at scale.

Cheaper intelligence should expand—not reduce—the artificial intelligence (AI) opportunity. Lower costs increase adoption. As AI becomes embedded across more workflows, usage can grow much faster than prices fall.

The winners won't depend on one model. Competitive advantage increasingly belongs to companies that can orchestrate many models, manage complexity and deploy intelligence efficiently.

Trust may become AI's most valuable product. Reliability, governance, security and accountability are becoming just as important as raw model performance.

Intelligence is getting cheaper

Open-weight AI is improving quickly. Models that once required enormous budgets can now be downloaded, adapted and deployed at far lower cost. Token prices are falling, performance gaps are narrowing and capabilities are diffusing faster.

This progress has produced a straightforward bearish argument: If intelligence becomes a commodity, the laboratories and data centers built to produce it will struggle to earn attractive returns. Lower prices will produce lower revenue, weaker margins and, eventually, less investment in AI infrastructure.

We think that conclusion is wrong.

Enterprises do not ultimately buy intelligence by the token. They buy work. They buy trusted outcomes. They want code that ships, claims that are underwritten, customer issues that are resolved, research that produces useful conclusions and decisions that withstand scrutiny.

The model matters. But it represents only one part of the system that produces those outcomes.

The right measure is cost per completed task

The relevant economic measure is not cost per token. It is the cost of completing a task reliably.

A cheaper model may require repeated attempts, additional verification, human supervision and correction. A more expensive model may complete the same task correctly on the first attempt.

Comparing the two solely by token price is like choosing a contractor based on hourly wages while ignoring how long the job will take, how much rework it will require and whether the finished result will meet the necessary standard.

Not all contractors are the same. Neither are all models, even when benchmarks make them appear comparable.

The more important question is not how cheaply a model can generate an answer. It is how efficiently the full system can produce a trustworthy outcome.

The model is just one part of the system

A model does not operate in isolation. To perform valuable work, it needs proprietary data, organizational context, tools, permissions, workflow logic, security controls, evaluations and clear rules for when a person must intervene.

The model supplies intelligence. The surrounding system determines whether an enterprise can use that intelligence safely, reliably and efficiently.

Think of this surrounding layer as the harness.

The harness selects the appropriate model, connects it to data and tools, evaluates its outputs, enforces permissions, logs its actions and determines when human escalation is required. Over time, it accumulates workflow logic, institutional memory, security policies and knowledge of where systems are most likely to fail.

An enterprise may replace the underlying model several times while leaving much of the harness in place.

That persistence matters. As models become increasingly interchangeable, the system around the model may become more valuable.

Exhibit 1: A model is powerful. The harness makes it enterprise-ready.

Value will migrate, not disappear

The commoditization of model access does not mean the economics of AI disappear. It means value should move toward the layers that remain scarce:

Infrastructure

Distribution

Proprietary context

Workflow ownership

Governance

Trust

The weakest business model may be selling undifferentiated access to a model while relying on a temporary performance advantage to sustain pricing.

Advances diffuse. Customers switch providers. Applications route work among different models. Open alternatives establish a lower-cost reference point.

Frontier model providers understand this threat. They are moving up the stack into enterprise platforms, agent systems, security, workflow software and, ultimately, completed work.

We believe that is the right strategy. But it will not be easy.

Entrenched platforms, software providers, service firms and enterprises already occupy the layers above the model. They own the data, workflows, distribution and customer relationships. They will defend those positions aggressively.

Moving up the stack also creates conflict. The further a model provider expands into legal work, software development, financial analysis, customer service or research, the more likely its customers are to view it as a competitor rather than a neutral supplier.

That tension should accelerate demand for portability, private deployment and systems that can use multiple models.

Falling prices do not necessarily reduce the revenue opportunity. They can expand the market."

Cheaper intelligence should increase consumption

Falling prices do not necessarily reduce the revenue opportunity. They can expand the market.

A traditional software query may involve one request and one response. An agentic workflow—one in which an AI system plans and completes a sequence of actions—can require dozens or hundreds of model calls.

The system may need to create a plan, retrieve information, use tools, test alternatives, correct mistakes and verify the final result.

As intelligence becomes cheaper, companies will apply it to more problems and automate more steps. They will use it to serve customers better, accelerate growth, improve decisions and operate more efficiently.

The cost of each unit of intelligence can fall while the number of units consumed rises much faster. This is Jevons

Paradox applied to cognition: When a resource becomes cheaper and more efficient to use, total consumption can rise rather than fall.

Exhibit 2: Lower AI Costs Appear To Be Driving More Usage, Not Reducing Revenue Opportunity

January 2024-April 2026

Source: Exponential View analysis; Epoch AI. Chart simplified from original to show token price and usage intensity only. Blended price per million tokens vs tokens processed per output token. Interpretations regarding AI adoption, usage, market expansion and commercial impacts reflect current market observations and expectations based on available industry data and are subject to change. There can be no assurance these trends will continue or that anticipated commercial outcomes will be realized.

Cheaper intelligence should therefore expand both the market for intelligence and the infrastructure required to produce it. The critical variable is not simply the decline in unit price. It is how strongly consumption responds.

The strongest systems will manage abundance

We believe the strongest position may belong to companies that can use many models rather than depend on one.

They can route each task to the provider offering the best combination of quality, reliability, latency, security and cost. They can reserve frontier intelligence for the hardest problems while assigning routine work to smaller and less expensive models.

The result will be disciplined abundance.

Enterprises learning how to manage that abundance today may hold a meaningful advantage tomorrow.

For investors, the key question is not simply who has the best model. It is who can convert abundant intelligence into completed work at the lowest total cost-and do so with enough reliability and trust that customers redesign their operations around the system.

Winning Systems Won't Rely on a Single Model

Infrastructure providers should benefit if usage grows faster than unit prices decline. Computing platforms can create value by aggregating models and managing deployment, security, governance and billing. Applications can build durable advantages when they own critical workflows, proprietary context and customer relationships.

Enterprises that deploy these systems effectively may capture a substantial share of the value. They can use AI to raise productivity, improve decisions, increase capacity and grow faster.

The weakest position is likely to be the undifferentiated middle: an expensive model or thin application with no sustainable cost advantage, proprietary data, distribution, workflow ownership or accountability for the outcome.

Trust may become the most valuable layer

For consequential work, customers will pay for reliability, explainability, security, auditability and clear responsibility when something goes wrong.

The cheapest model is irrelevant if an organization cannot approve it for production.

Open-weight models may commoditize access to intelligence. They will not commoditize the ability to apply intelligence reliably inside a complex enterprise. Nor will they reduce demand for intelligence. We believe they will expand it meaningfully.

Compute is not free. Workflow is not free. Proprietary context is not free. Governance is not free. Accountability is not free.

The winners may not have the smartest model at every moment or for every task. They will be the companies that turn abundant intelligence into trustworthy, completed work.

They will embed themselves in how enterprises operate. They will create value by helping customers serve clients better, grow faster and make better decisions.

Exhibit 3: Successful Enterprise AI Outcomes

The Bottom Line

The prize is large. The noise will be loud. The volatility will likely be high as the market determines where the value ultimately settles.

We believe investors who navigate this transition successfully will need to be nimble, active and focused on the central idea:

The supply of intelligence is increasing rapidly. The companies that can translate it efficiently into trusted, explainable work will capture the greatest value.

Endnote

Open-weight AI refers to AI models whose trained parameters—the numerical “weights” that encode what the model learned—are made available for others to download and run.

WHAT ARE THE RISKS?

All investments involve risks, including possible loss of principal.

Equity securities are subject to price fluctuation and possible loss of principal.

Small- and mid-cap stocks involve greater risks and volatility than large-cap stocks.

Investment strategies which incorporate the identification of thematic investment opportunities, and their performance, may be negatively impacted if the investment manager does not correctly identify such opportunities or if the theme develops in an unexpected manner. Focusing investments in the health care, information technology (IT) and/or technology-related industries carries much greater risks of adverse developments and price movements in such industries than a strategy that invests in a wider variety of industries.

Any companies and/or case studies referenced herein are used solely for illustrative purposes; any investment may or may not be currently held by any portfolio advised by Franklin Templeton. The information provided is not a recommendation or individual investment advice for any particular security, strategy, or investment product and is not an indication of the trading intent of any Franklin Templeton managed portfolio.

WF: 11749558